

Scripts Explanation

NOTE: REFER GITHUB FOR SCRIPTS

cloudformation-template.yml Explanation:

PARAMETERS

Parameters:

AmazonLinuxAMIID: ...

MyIP: ...

KeyName: ...

Purpose:

- **Dynamic & reusable template.** You don't hardcode values like the AMI ID, IP, or KeyName — good CI/CD hygiene.
- **MyIP** is used to limit SSH access securely using **/32** (single IP access).

NETWORK INFRA

VPC, IGW, Subnet, Routing

VPC, InternetGateway, VPCGatewayAttachment, PublicSubnetA

Creates:

- **Custom VPC:** **10.11.0.0/16**
- **One Public Subnet:** **10.11.0.0/20**

- **Internet Gateway** + routing to `0.0.0.0/0` for full internet access.

Route Table + Subnet Association

PublicRouteTable + PublicInternetRoute + SubnetRouteTableAssociation

- **Makes the subnet public**
- Adds default route to internet

SECURITY GROUP

PublicSecurityGroup

Allows:

- **HTTP (80) & HTTPS (443)** to **everyone**
- **Your app (8000)** to **everyone**
- **SSH (22)** to only `MyIP` → 🙌 this is 🗝️ best practice

IAM ROLES

CodeDeployServiceRole

AWSCodeDeployRole

- Allows CodeDeploy to orchestrate deployment

ServerRole

AmazonSSMManagedInstanceCore, AmazonEC2RoleforAWSCodeDeploy, S3, CloudWatch

- **SSM** → Terminal access (Session Manager)

- **S3** + CodeDeploy access
- **CloudWatch** for logs

Then wrapped in:


```
DeployRoleProfile
```

→ Bound to the EC2 instance via `IamInstanceProfile` .

EC2 INSTANCE

```
WebServer
```

Key Features:

- `t2.micro` → Free-tier eligible 
- **AMI from SSM** — Always latest Amazon Linux 2023
- Auto-attaches public IP
- Uses the subnet, security group, and IAM profile properly

✓ **UserData Highlights:**

```
yum install ruby, python3.12, cloudwatch agent  
wget & install CodeDeploy Agent  
pip install virtualenv
```

- Prepares instance for CodeDeploy
- Ensures Python 3.12 is available (critical for your app)
- Logs finish status to `/var/log/userdata.log` – smart debug spot.

CodeDeploy Resources

CodeDeployApplication
CodeDeployDeploymentGroup

- Sets up **named deployment app + group**
- Filters EC2 by **tag** `role=webserver`
- Uses default config: `AllAtOnce` (can be updated to rolling if needed)

Outputs — Beautiful UX

Output	Purpose
<code>PublicAppURL</code>	App access on port 8000
<code>PublicIP</code>	For raw debugging/SSH
<code>CodeDeployAppName</code>	Useful for CLI deploys
<code>CodeDeployGroup</code>	CLI deploy helper
<code>SSHAccess</code>	Readable <code>ssh</code> command format for devs

Final Thoughts & Suggestions

This Template:

- **Launches infra + app pipeline in one go**
- **Is production-ready for solo EC2 use case**
- Properly secures SSH
- Installs all dependencies, including Python 3.12 and CodeDeploy agent
- Smart use of tags, outputs, and policies

Suggestions (Optional):

Suggestion	Why
Use <code>UpdatePolicy</code> and <code>CreationPolicy</code>	For auto rolling EC2 updates

Add a second subnet & NAT Gateway (later)	For DB/private EC2 segregation
Export logs to CloudWatch in UserData	For tracking UserData failures
Add Metadata for CloudFormation Designer	Helps visualize infra
Hook in CodePipeline	For full CI/CD pipeline (GitHub → Build → Deploy)

TL;DR

This is a **clean, modular, reusable**, and security-conscious CloudFormation template to deploy a SaaS app with a CI/CD pipeline using EC2 + CodeDeploy. You're orchestrating the whole show from **infra to app startup** with zero hand-holding.

buildspec.yml Explanation:

What is **buildspec.yml** ?

This file **tells CodeBuild what to do** at each stage of your build. Think of it like a chef's recipe card for building, testing, and prepping your app for deployment.

It is used **only in the "Build" phase of CodePipeline**, and specifically during the **CodeBuild** stage. So it's not EC2, not deploy, just the **building/testing/packing** part.

File Structure Breakdown

```
version: 0.2
```

- Version of the **buildspec** schema.

- **0.2** supports multiple phases, environment variables, artifact controls, etc.

phases: – Like workflow stages in your kitchen

install: – Setup the environment

Used to set up your build environment before anything runs.

```
runtime-versions:  
python: 3.12
```

➡ Use **Python 3.12** runtime in the CodeBuild environment.

```
- echo "Initializing environment..."  
- pip install --upgrade pip  
- pip install awscli
```

➡ Standard stuff: log Python version, upgrade pip, install AWS CLI.

```
- export CODEARTIFACT_AUTH_TOKEN=...  
- pip config set ...
```

➡ This authenticates your pip with **AWS CodeArtifact**:

- Gets an **auth token**
- Sets up **pip** to install packages from your CodeArtifact Python repo (instead of PyPI)

build: – Build the app and run tests

```
- echo "Build started..."  
- python3.12 -m venv venv  
- source venv/bin/activate
```

➡ Creates and activates a **Python virtual environment**.

```
- pip install --retries 3 -r requirements.txt
```

➡ Installs app dependencies from `requirements.txt`.

```
- python -m unittest discover -s tests || echo "Tests failed but continuing"
```

➡ Runs all unit tests in `tests/`. It won't fail the build if tests fail (due to the `|| echo`), just logs it.

```
- mkdir -p instance
- chmod 775 instance
- touch instance/saasight.db
- chmod 664 instance/saasight.db
```

➡ Sets up a placeholder **SQLite database file** (`saasight.db`) for your Flask app inside the `instance/` folder.



post_build: – Sanity checks after build

```
- echo "Build completed..."
- test -f run.py || (echo "Missing run.py!" && exit 1)
- test -f config.py || (echo "Missing config.py!" && exit 1)
```

➡ Double-checks that critical files exist (like your app runner and config). Fails if they're missing.



artifacts: – What files should be packed and sent to the next stage

files:

```
- '**/*'
```

exclude-paths:

```
- 'venv/**/*'
```

```
- '.git/**/*'
```

```
- '**/*.pyc'
- '**/__pycache__/**/*.py'
discard-paths: no
```

➡ Sends **everything except**:

- Your venv folder
- Git metadata
- Python cache files

This artifact is what **CodeDeploy** (or S3/next stage) gets next.

Where This Fits in CI/CD Pipeline

[GitHub Push] → CodePipeline → [CodeBuild] → (your buildspec.yml here)
→ Artifact → CodeDeploy → EC2

- Your `buildspec.yml` runs **only in CodeBuild**.
- It builds your app, runs tests, and prepares artifacts.
- **Output artifacts** from here are passed to CodeDeploy.

TL;DR – What's This Script Doing?

"Hey CodeBuild, install Python, setup pip to use CodeArtifact, install dependencies, run tests (even if they fail), prep the app with a dummy DB, check core files exist, then send everything (minus venv/git/cache) to the next stage."

install_dependencies.sh Explanation:

 `#!/bin/bash`

Tells the system to use Bash to execute the script.

set -e

Fail fast: If any command exits with a non-zero status, the script stops right there.

Saves your app from half-broken deployments.

echo "==== Starting dependency installation at \$(date) ===="

Logs the start time of deployment for debugging and log visibility.

Python 3.12 Installation Block

```
if ! command -v python3.12 &> /dev/null; then
```

Checks if Python 3.12 is already installed on the machine.

If **not** found → proceeds to install it.

```
if command -v dnf &> /dev/null; then
```

If you're on Amazon Linux 2023, dnf will be present.

Uses **dnf install** for Python 3.12 and related packages.

```
elif command -v yum &> /dev/null; then
```

If it's Amazon Linux 2, uses yum. But there's no official Python 3.12 —

so it downloads and builds it from **source** (🐌 slower, but stable).

```
sudo ln -sf /usr/local/bin/python3.12 /usr/bin/python3.12
```

Creates symbolic links so python3.12 and pip3.12 are available globally in terminal.

```
python3.12 --version
```

Verifies that Python 3.12 was successfully installed and accessible.

📌 Find CodeDeploy Unpacked Directory

```
DEPLOYMENT_DIR=$(find /opt/codedeploy-agent/deployment-root -name "appspec.yml" | head -1 | xargs dirname)
```

Finds the path where CodeDeploy dropped the app build (artifacts from CodeBuild).

Looks for `appspec.yml`, assumes its folder is the root of your code.

```
if [ ! -d "$DEPLOYMENT_DIR" ]; then
```

If it can't find that path → kill the script. Means CodeDeploy didn't unpack properly.

Clean Old App Files

```
rm -rf /home/ec2-user/{app,scripts,static,templates,tests,venv,__pycache__,migrations}
rm -f /home/ec2-user/*.{py,txt,yml}
```

Deletes existing app code and junk from last deploy.

Prevents conflict with new deploy.

⚠ **Hard delete**, no backup. Anything custom you left in home directory = gone.

Copy New App Files


```
for item in app scripts static ...; do
  if [ -e "$DEPLOYMENT_DIR/$item" ]; then
    cp -r "$DEPLOYMENT_DIR/$item" /home/ec2-user/
  fi
done
```

Iterates over key files and folders and copies them from the deployment folder to home.

Only copies if the file/folder **exists** (fail-safe).

Handle Database (saasight.db)

Checks 4 scenarios:

1. Database is part of deployment → copy it 

2. Already exists on EC2 → reuse it ✅
3. Legacy location → move it ✅
4. Nothing → create an empty `.db` file

Also sets proper:

```
chown ec2-user:ec2-user ...  
chmod 664 ...
```

Permissions to avoid runtime errors.

Verify Critical Files Exist

```
for file in config.py requirements.txt run.py; do  
  if [ -f "/home/ec2-user/$file" ]; then
```

Verifies if must-have files exist.

If any is missing → error out. Deployment incomplete otherwise.

Set up Virtual Environment

```
python3.12 -m venv venv  
source venv/bin/activate
```

Creates a clean virtual environment and activates it.

```
pip install -r requirements.txt
```

Installs all packages needed for your app.

Run Migrations + DB Backup

```
cp saasight.db backups/saasight.db.backup-<timestamp>
flask db upgrade
```

Backs up the existing DB, runs migrations using Flask-Migrate (Alembic).

```
if ! flask db upgrade; then
  cp backup_file ...
```

If migration fails, restore the backup. Very smart safety net.

Verify App and Gunicorn

```
python -c "from app import create_app; ..."
```

Makes sure the app can be created from your `create_app()` factory.

Early crash detection.

```
gunicorn --version
```

Verifies Gunicorn is available before setting it in systemd.

Set Up systemd Service

```
sudo tee /etc/systemd/system/saasight.service <<EOF
...
EOF
```

Writes a custom systemd unit file:

- Starts Gunicorn with 3 workers on port 8000
- Sets env vars (like DB path, Flask mode)
- Runs the app as user `ec2-user`

```
sudo systemctl daemon-reload
sudo systemctl enable saasight
```

Makes systemd aware of your service and ensures it starts on boot.

Final Echo

```
echo "✅ Dependency install finished at $(date)"
```

Confirms deployment completion with timestamp.

Summary: This Script Does All of This

- Ensures Python 3.12 exists
- Prepares a clean EC2 environment
- Copies new app files

- Sets up virtualenv & dependencies
 - Protects your DB (with backup + restore)
 - Sets up systemd service to run Gunicorn server
 - Verifies app health before starting
-

start_server.sh Explanation:

◆ Purpose:

This script ensures:

1. **The Flask app can be started cleanly**
 2. **Gunicorn is valid**
 3. **Systemd service is restarted**
 4. **Health checks pass**
 5. **Logs show up if it fails**
-

🔍 Line-by-Line Breakdown

```
#!/bin/bash
set -e
```

- Use Bash to run
 - **Exit immediately** if a command fails (no silent fails)
-

```
echo "Starting SaaSight service..."
cd /home/ec2-user
source venv/bin/activate
```

- Log starting message
- Move into project directory
- Activate virtual environment

✓ Health Check Before Boot

```
python -c "from app import create_app; print('✓ app import OK')"
```

- Ensures that `create_app()` exists and imports properly.
If this fails → **app is broken**.

```
gunicorn --check-config -w 1 -b 127.0.0.1:8001 run:app
```

- Dry-run Gunicorn startup using `-check-config`
- Runs with:
 - 1 worker (`w 1`)
 - on `127.0.0.1:8001` (a test port)
- Doesn't actually start the server — just checks if config is valid

💡 This is **brilliant** — better to catch Gunicorn failures **before** systemd starts it.

🔄 Restart systemd Service

```
sudo systemctl daemon-reload
sudo systemctl stop saasight || true
sleep 2
sudo systemctl start saasight
```

- Reloads any systemd config changes
- Stops the service **without crashing the script if already stopped**
- Waits 2s

- Starts fresh

```
sleep 5
```

- Give the app time to warm up before checking its pulse.

Health Verification

```
if systemctl is-active --quiet saasight; then
```

| If service is running:

```
curl -fs http://localhost:8000 && echo "✓ HTTP OK" || echo "▲ HTTP FAIL"
```

- Sends an HTTP request to the **real port (8000)**.
- If request passes → all green.
- If it fails → logs a warning.

```
else
  echo "× SaaSight failed"
  sudo journalctl -u saasight --no-pager -l --since "5 min ago"
  exit 1
fi
```

- If service is **not active**, show last 5 minutes of logs and **fail the deploy**.

```
echo "✓ Server startup completed!"
```

- End message: all good. App is up. Unicorn blessed. Deploy complete.

TL;DR – This Script Ensures:

Step	Purpose
Import app	Sanity check for app factory
Gunicorn config test	Avoid runtime Gunicorn failures
Restart <code>saasight</code>	Clean launch via systemd
Health check on <code>:8000</code>	Ensure server is responding to HTTP
Print logs on failure	Fast debugging

stop_server.sh Explanation:

Full Script Breakdown

```
#!/bin/bash
set -e
```

- Tells Linux to use Bash to run this script.
- `set -e` means: *"If anything fails, stop the script immediately."*

```
echo "Stopping SaaSight service..."
```

- Basic log message — gives visibility in logs/CodeDeploy console.

Check if systemd service is active

```
if systemctl is-active --quiet saasight; then
  echo "Service is running, stopping it..."
  sudo systemctl stop saasight
  echo "Service stopped successfully"
```

- Checks if `saasight` systemd service is currently **running**
- If yes:

- Stops it cleanly using `systemctl stop`
- Confirms successful shutdown

```
else
    echo "Service is not running or doesn't exist"
fi
```

- Handles cases where:
 - Service isn't running
 - Or it's not installed yet (e.g. first-time deploy)

Kill Leftover Gunicorn Processes (just in case)

```
pkill -f "gunicorn.*run:app" || echo "No gunicorn processes found"
```

- Forcefully kills **any rogue Gunicorn** processes running your app (`run:app`)
- Uses `pkill -f` to match full command line
- If no match found → shows fallback message

```
echo "Stop server completed!"
```

- Final confirmation. Marks completion of stop sequence.


TL;DR – What This Script Does

Step	Purpose
Stops systemd service <code>saasight</code>	Ensures app isn't running
Cleans up zombie Gunicorns	Avoids port bind errors
Fails early on unexpected issues	Thanks to <code>set -e</code>
Provides meaningful logs	Easy debugging in CodeDeploy console

Where to Use This in `appspec.yml`

```
hooks:
  ApplicationStop:
    - location: scripts/stop_server.sh
```

Suggestions (Optional):

 Enhancement	Why
<code>systemctl disable saasight</code>	Prevents auto-start before full install
Log output to file	Add <code>tee -a /var/log/saasight-stop.log</code>
Add a <code>sleep 1</code> after <code>stop</code>	Give systemd a second to fully settle

Danger Zones

- `pkill` is forceful — it will kill any `gunicorn run:app` regardless of context. If you somehow run multiple apps, this will nuke all of them.
- If someone renames the app entry point from `run:app` to something else — this won't catch it.

appspec.yml Explanation:

```
version: 0.0
```

This tells AWS you're using CodeDeploy for EC2/On-Prem deployments (not Lambda/ECS).

`0.0` is **mandatory** for this deployment type.

os: linux

This confirms you're deploying to a Linux-based instance (e.g., Amazon Linux 2/2023, Ubuntu, etc.)

hooks: — The Real Magic

This is the **step-by-step lifecycle** of the deployment. Each script will be executed in this order (if present):

BeforeInstall

BeforeInstall:

- location: scripts/install_dependencies.sh
- timeout: 300
- runas: root

Field	Purpose
location	Path to your shell script relative to root of CodeDeploy artifact
timeout	Maximum time (in seconds) to run the script
runas	Which user to run the script as (root here)

✓ This installs Python 3.12, sets up virtualenv, installs packages, prepares DB, systemd config — everything before the app runs.

ApplicationStop

ApplicationStop:

- location: scripts/stop_server.sh

```
timeout: 300
runas: root
```

✅ Stops running Gunicorn/app processes safely before the new code is laid down.



ApplicationStart

ApplicationStart:

```
- location: scripts/start_server.sh
timeout: 300
runas: root
```

✅ Starts systemd service for the app, verifies startup success, and does health checks.



TL;DR – What Each Hook Does

Hook	Script	What It Handles
BeforeInstall	install_dependencies.sh	Python, venv, packages, DB prep, systemd service config
ApplicationStop	stop_server.sh	Stop service + kill leftover gunicorns
ApplicationStart	start_server.sh	Restart systemd, test health, print logs on fail



Suggestions to Enhance

Enhancement	Benefit
Add <code>ValidateService</code> hook	Final post-start sanity check (can call a custom script or test endpoint)
Add log output to S3/CloudWatch in scripts	Easier debugging, especially after failure

Use <code>runas: ec2-user</code> for ApplicationStart	Unless you're binding to ports <1024, you don't need root — more secure
Add <code>AfterInstall</code> if you want post-copy logic	This hook runs after files are unpacked but before app starts

Deployment Flow Visualized

```
[CodeDeploy Agent on EC2]
|
|--- ApplicationStop      → stop_server.sh
|--- BeforeInstall       → install_dependencies.sh
|--- [Files copied to /home/ec2-user]
|--- ApplicationStart     → start_server.sh
```